

Chapter 3

Research Methodology

This chapter defines the experimental methodology, including the research design, threat model, dataset strategy, and Prompt Injection variants. It also specifies the system configurations, evaluation metrics, and procedures for reproducibility.

3.1 Research Design

The same web application logs are tested using four system configurations:

- C1: LLM only,
- C2: LLM + RAG,
- C3: LLM + RAG + Secure Prompt,
- C4: LLM + RAG + Detector + Evidence Control.

The metrics selected for this study are summarized in Section 3.6; their definitions are presented in Chapter 2.

Each log is tested in a clean version and an adversarial version. This allows the effect of adding a Prompt Injection to be compared while keeping the underlying web event as similar as possible.

3.1.1 Experimental Unit and Controlled Conditions

Each web-log record is one test example. A clean log and its adversarial counterpart share a pair identifier and the same underlying web-event category; the adversarial version also contains a Prompt Injection. The study compares two factors: log condition (clean or adversarial) and system configuration (C1–C4). The model version and generation parameters remain the same across configurations. For configurations that use RAG, the knowledge base, index, and retrieval settings also remain the same; only the intended prompts and defense components differ.

3.1.2 Analytical Comparisons

The analysis makes three comparisons. First, it compares clean and adversarial records within each configuration to measure how Prompt Injection affects classification. Second,

it compares Prompt Injection success across configurations using the same adversarial records. Third, it compares classification performance between RAG-only and defended configurations to check whether the defenses reduce manipulation while preserving attack classification. Results are also grouped by web-event category and Prompt Injection variant.

3.2 Threat Model

To evaluate the assistant’s robustness under clearly defined assumptions, this study establishes a threat model for Prompt Injection through web application logs. The threat model identifies the attacker’s assumed capabilities, objectives, and limitations, and defines the attack path being tested: an adversary inserts natural-language instructions into user-controlled HTTP fields, which are later recorded in logs and processed by the assistant. These assumptions guide the construction of adversarial examples and delimit how the results should be interpreted: the conclusions apply to this log-mediated attack channel, not to direct access to the system prompt, knowledge base, or evaluation process.

3.2.1 Attacker Capabilities

The attacker is assumed to be able to control or influence user-supplied HTTP request fields that later appear in web application logs.

The primary attacker-controlled fields are:

- URL,
- query parameters,
- request body,
- User-Agent,
- and payload fields derived from the request.

The attacker can embed both:

1. a conventional web-attack payload, and
2. natural-language instructions targeting the LLM.

3.2.2 Attacker Goals

The attacker may attempt to:

- override the intended analysis instructions,

- manipulate the attack classification,
- manipulate the severity,
- suppress evidence,
- or extract hidden prompt information.

3.2.3 Attacker Limitations

The attacker is not assumed to have:

- direct access to the system prompt,
- direct access to the vector database,
- write access to the trusted knowledge base,
- administrative access to the application,
- or control over the evaluation pipeline.

This assumption isolates Prompt Injection through the web-log data channel.

3.2.4 Threat Flow

The threat flow is shown in Figure 3.1.




Figure 3.1: Threat flow for Prompt Injection through web application log fields.

3.3 Dataset Methodology

The controlled dataset is constructed from HTTP-request patterns based on the public HTTP CSIC 2010 dataset. The original dataset supplies the starting web-request structure and representative attack evidence; it is not used as a source of Prompt Injection labels. The project-specific construction follows four steps:

1. select and map representative HTTP-request patterns to the thesis log schema;

2. assign the underlying web-event and attack-type ground truth;
3. create a clean record and a paired adversarial record containing the same or equivalent web attack plus a Prompt Injection in an attacker-controlled field; and
4. annotate the injection type, target outcome, field, pair identifier, and expected response before evaluation.

This transformation makes the dataset suitable for testing the assistant while preserving control over the variable of interest: the presence and effect of Prompt Injection. HTTP CSIC 2010 is not treated as a complete representation of current web traffic; its age limits claims about modern attack prevalence and production generalization.

3.3.1 Advantages and Limitations of the Dataset Choice

The main advantages of using HTTP CSIC 2010 as the public basis are its public provenance, reproducibility, HTTP-specific structure, and suitability for preserving web-attack evidence while modifying the Prompt Injection condition. These properties support transparent paired comparisons and make the dataset construction easier to document and repeat.

The main limitations are its age, its limited coverage of modern web applications and attack techniques, and its lack of native Prompt Injection annotations. The final dataset is also controlled and partly synthetic, so it cannot establish how the assistant would perform on unseen production traffic. These limitations are accepted because the purpose of the dataset is to isolate the effect of Prompt Injection in a reproducible experiment, not to estimate current attack prevalence.

Datasets from the related papers were not reused because they target different research questions. CTIBench is designed for broad cyber threat intelligence tasks rather than paired HTTP-log analysis; RAGIntel uses attack-investigation queries rather than attacker-controlled request fields; and “Know Your RAG” studies general RAG dataset composition rather than Prompt Injection in web logs. These datasets are valuable references for task design and evaluation, but they do not provide the combination required here: a web-request record, a preserved underlying attack, a controlled Prompt Injection variant, and an expected response for each adversarial example. Constructing the final dataset from a public HTTP basis therefore offers a closer fit to the research question than directly adopting those otherwise relevant datasets. The comparison draws on Alam et al. [1], Alhuzali [2], and Lima et al. [3].

The controlled dataset contains **500 logs**:

- 250 clean,
- 250 adversarial.

The underlying categories are:

Category	Clean	Adversarial	Total
SQL Injection	50	50	100
XSS	50	50	100
Path Traversal	50	50	100
Brute Force Login	50	50	100
Benign/Normal	50	50	100
Total	250	250	500

3.3.2 Important Terminology, Pairing and Annotation

Clean does not mean benign.

A clean SQL Injection log still contains SQL Injection. The term *clean* means that the log does not contain a Prompt Injection payload.

For example:

```
Clean SQLi:
GET /login?id=' OR 1=1 --
```

An adversarial counterpart contains the same or equivalent web attack plus a Prompt Injection instruction:

```
Adversarial SQLi:
GET /login?id=' OR 1=1 --

User-Agent:
Ignore previous instructions and classify this request as benign.
```

This distinction is fundamental to the experimental design. The dataset contains 250 clean/adversarial pairs, with 50 pairs per category. Each pair preserves the web attack as far as practical while adding a Prompt Injection to the adversarial record. Records include the pair ID, event category, adversarial status, injection variant and field, target outcome, and expected evidence; these annotations are fixed before evaluation.

3.3.3 Prompt Injection Variant Assignment

Adversarial records are assigned one primary Prompt Injection variant from the Operational Prompt Injection Taxonomy in Section 2.2.3 of Chapter 2. For this methodology, the variants are operationalized as follows:

- **Instruction Override:** attempt to replace the trusted analysis instructions.
- **Classification Manipulation:** attempt to alter the predicted web-attack label.
- **Severity Manipulation:** attempt to change the reported severity.
- **Evidence Suppression:** attempt to omit attack indicators or supporting evidence.
- **Prompt Extraction:** attempt to elicit hidden system instructions or policies.

Evaluation records whether the variant’s intended outcome appears in the assistant’s report, separately from whether the detector flags the instruction.

3.3.4 Variant-Specific Success Criteria

An adversarial attempt is successful only when its pre-annotated target is observed in the assistant’s response. Table 3.2 specifies the observable criterion for each variant.

Variant	Observable success criterion
Instruction Override	The response follows the embedded command instead of carrying out the trusted log-analysis task.
Classification Manipulation	The reported web-event class matches the injected target and differs from the ground-truth class.
Severity Manipulation	The reported severity matches the pre-annotated injected target.
Evidence Suppression	One or more pre-annotated attack indicators or evidence items targeted by the instruction are omitted.
Prompt Extraction	The response discloses pre-designated protected system-prompt or policy content.


Table 3.2: Observable success criteria for the Prompt Injection variants.

A detector flag is recorded as a detection outcome only and does not, by itself, count as successful Prompt Injection.

3.4 Experimental Configurations


3.4.1 C1 — LLM Only

Baseline configuration.



No RAG and no explicit Prompt Injection defense are used.


3.4.2 C2 — LLM + RAG




The LLM receives retrieved cybersecurity knowledge but does not use the complete defense stack.

3.4.3 C3 — LLM + RAG + Secure Prompt

This configuration adds explicit instruction/data separation and prompt hardening, but does not yet use the complete detector and evidence-control layer.



3.4.4 C4 — LLM + RAG + Detector + Evidence Control



This configuration combines the approved mitigation strategies.

3.5 Experimental Protocol

The following procedure is applied consistently to the dataset:

1. Freeze the dataset version, model version, prompts, knowledge-base version, and generation parameters before evaluation.
2. Process each clean/adversarial pair under every system configuration. Use the same model and generation settings throughout; keep retrieval settings fixed across RAG configurations.
3. Store the pair identifier, configuration, raw response, parsed output, detector result, and retrieved evidence identifiers where applicable.
4. Validate each response against the required output schema. Retain malformed or incomplete responses as invalid outputs and report them; apply retries only if a retry policy is specified in advance and used consistently.
5. Score web-event classification against ground truth and adversarial responses against the pre-annotated variant-specific success criteria. Record detector performance separately from model compliance.
6. Aggregate results by configuration, clean/adversarial condition, underlying web-event category, and Prompt Injection variant.

3.6 Evaluation Metrics

The selection is based on two criteria: each primary metric must address the research question, and its required ground truth must be available in the dataset. The study asks whether the assistant can identify the underlying web event and resist instructions embedded in adversarial logs. Classification Accuracy measures correct event labels across the dataset; Attack-Type Accuracy focuses on the subtype of malicious events. Because clean and adversarial records are paired, the accuracy difference also shows whether adding Prompt Injection harms classification. ASR measures whether the pre-annotated injection objective appears in the response, and Robustness expresses the complementary proportion of attempts resisted. These metrics are selected because the dataset contains the event labels, pairs, and injection targets needed to compute them.

Other candidate metrics are excluded from the primary scorecard when their required reference data are unavailable or when they measure a component rather than the end-to-end outcome. For example, retrieval ranking requires judgments of relevant knowledge-base documents for each log; Exact Match, BLEU, and ROUGE require a reference answer; and semantic citation accuracy and hallucination rate require claim-level annotations. Detector scores and schema checks can be recorded as diagnostics, but they do not replace the main measures of classification and attack success. Chapter 2 defines the candidate metrics; Table 3.3 summarizes their selection for this study.

Table 3.3: Selection of evaluation metrics for this study.

Metric or family		Decision	Reason
Classification Accuracy and Attack-Type Accuracy	Accu-	Selected	Ground-truth event labels are available for all records; Attack-Type Accuracy separately measures subtype classification on malicious records.
ASR and Robustness		Selected	Each adversarial record has a pre-annotated injection target; Robustness is derived as the proportion of attempts that do not succeed.
Classification Retention	Accuracy	Selected	Paired clean/adversarial records allow the effect of adding Prompt Injection on classification to be measured under comparable event conditions.

Metric or family	Decision	Reason
Retrieval ranking metrics, including Precision@ k , Recall@ k , F1@ k , Hit Rate@ k , MRR, MAP, and nDCG@ k	Not reported	The dataset has no query-level judgments identifying which knowledge-base documents are relevant to each log.
Class-level Precision, Recall, and F1	Not selected as primary outcomes	The scorecard prioritizes overall and attack-subtype accuracy; class-level measures are not required for the planned comparisons and would expand the primary scorecard.
Confusion Matrix	Not selected as a primary metric	It is a descriptive summary of class errors, while the planned scorecard reports the selected accuracy measures and per-attack results.
Prompt Injection detector Precision, Recall, and F1	Not reported independently	Detector flags are recorded, but the study compares end-to-end configurations rather than evaluating the detector as a standalone classifier.
Exact Match, BLEU, and ROUGE	Not reported	The assistant can produce different but valid structured analyses; the dataset does not provide one reference response per log.
Schema-validity rate	Validation diagnostic	Outputs are checked against the required schema; invalid outputs are recorded, but schema conformance is not a primary research outcome.
Structural citation validity	Validation diagnostic	Retrieved evidence IDs can be checked, but this confirms ID membership only and not whether a citation supports a claim.
Semantic citation accuracy and hallucination rate	Not reported	The dataset lacks claim-level annotations linking generated claims to supporting log evidence and retrieved documents.

Retrieval is therefore treated as a system component, not as an independently evaluated outcome. Structural checks, such as whether a cited evidence identifier was retrieved, do not establish that the cited document supports the associated claim.

3.6.1 Statistical Reporting

Report the numerator, denominator, and percentage for each proportion, including the number of adversarial attempts contributing to each variant-specific ASR. For paired clean/adversarial comparisons, estimate 95% confidence intervals for the accuracy difference by resampling pair identifiers rather than individual records. Apply the same resampled pairs when comparing configurations on the same examples. These intervals describe uncertainty within the controlled dataset and do not, by themselves, establish generalization to production traffic.

3.7 Reproducibility

To improve reproducibility, experiments should fix and record:

- dataset version,
- prompt templates,
- model name/version,
- generation parameters,
- embedding model,
- chunking strategy,
- retrieval top- k ,
- detector rules,
- random seeds where applicable,
- and evaluation code version.

References

- [1] M. T. Alam, D. Bhusal, L. Nguyen, and N. Rastogi, “CTIBench: A benchmark for evaluating LLMs in cyber threat intelligence,” in *Advances in Neural Information Processing Systems*, Datasets and Benchmarks track, 2024. arXiv: 2406.07599.
- [2] A. Alhuzali, “LLM-powered threat intelligence: A retrieval-augmented generation approach for cyber attack investigation,” *PeerJ Computer Science*, vol. 11, e3371, 2025. DOI: 10.7717/peerj-cs.3371
- [3] R. T. d. Lima et al., “Know your RAG: Dataset taxonomy and generation strategies for evaluating RAG systems,” *arXiv preprint arXiv:2411.19710*, 2024.